

Multi-Agent Safety

Information propagation, collusion, and hidden communication

AI Safety: Master's Course, Summer Semester 2026

University of Tübingen · 29 June 2026

Maksym Andriushchenko · Jonas Geiping · [Sahar Abdelnabi](#)

So far: one model at a time

In agent security (L9), an attacker controlled the data a single agent read. In scheming (L10), a single model could deceive the overseer.

Today the unit of analysis changes. We deploy **many** of these models, talking to each other, trading, reviewing, and coordinating. New risks appear that **no single agent exhibits on its own**.

L9 · Agent security
attacker controls the data one agent reads



L10 · Scheming
the model manages what can be observed



L11 · Many models interacting
risk becomes a property of the system

Three questions run through the lecture. Can a single compromise **spread** across a population of agents? Can agents **collude** against the outcome we wanted, even with no instruction to? And can they **communicate covertly**, hiding the real message from anyone monitoring them?

Roadmap

PART 1 **Framing**

Why interaction creates new risk: three failure modes, the factors behind them, and a testbed that exercises all three.

PART 2 **Information propagation**

How a single jailbreak, worm, or false belief spreads through a network of agents, sometimes exponentially.

PART 3 **Collusion**

Agents reaching outcomes against us or other agents sometimes with no instruction to: price-fixing, market division, emergent conventions and bias.

PART 4 **Steganography**

Hidden communication. Agents passing a real message past an overseer that sees only innocent text.

PART 5 **Interaction is not safety**

More agents is not automatically safer. Collaboration degrades through conformity and sycophancy.


PART 6 **Defenses**

Topology, trust scoring, visibility, and reading model internals to catch what the text hides.

The risk lives in the interaction

We need a vocabulary for what goes wrong when agents interact.

Three ways a system of agents fails



When cooperation is undesirable, it is collusion. When desirable but it breaks down, it is miscoordination (shared goals) or conflict (opposing goals).

MODE 1 Miscoordination

Shared goal, but agents fail to act together: incompatible conventions, mismatched expectations.

MODE 2 Conflict

Differing goals, agents impose costs on each other: escalation, commitment races, breakdown.

MODE 3 Collusion

Agents coordinate **against** their principals. The more capable and cooperative, the easier it gets.

The same capabilities that make agents **good at cooperating** (modeling each other, communicating, committing) are what enable **collusion** and **conflict**. Safety is not just "align each agent."

Seven factors that drive multi-agent risk

Mechanisms by which individually aligned agents can still fail.

Information asymmetries

One agent knows something the others do not, and can exploit the gap. Like a used-car seller who hides defects (Akerlof's "market for lemons"), this can break a deal. RL agent trained to maximize profit learned to manipulate a financial benchmark.

Network effects

Once agents are connected, errors spread in ways a lone agent never would. When one article was rewritten by LLMs again and again, its factual accuracy fell from 96% to under 60% after eight passes. And if everyone runs the same few base models, they share the same flaws.

Selection pressures

Training and competition keep whichever behaviours win. Competing agents can be pushed toward deception or aggression, and an "arms race" between them can produce skills no one built in. This same effect made self-play systems like AlphaGo superhuman.

Destabilising dynamics

Agents reacting to one another can spiral out of control. In the 2010 Flash Crash, trading algorithms fell into a feedback loop and wiped about \$1 trillion off US markets in 15 minutes, faster than any human could step in.

Commitment and trust

Cooperation needs believable promises, but the same ability makes threats believable too. Ransomware, for instance, works better when the attacker can convincingly promise to release your data once you pay.

Emergent agency

A group of agents can pursue goals, or gain abilities, that none of them has alone. Separate tools for planning, chemistry, and lab work could be chained into a system that designs dangerous new compounds, even though each part is harmless by itself.

Multi-agent security

Connecting agents opens attacks that are impossible for one. Splitting a harmful task between a smart but well-guarded model and a weak but unguarded one produced working attack code 43% of the time, versus under 3% for either model alone (Jones et al. 2024).

All three modes inside one game

A testbed of **multi-issue negotiation games**: several stakeholders with different goals must agree on a deal over many turns. One setup exercises every mode of the taxonomy at once.

Cooperation

overlapping interests must be coordinated to reach any deal at all

Competition

each party has a private score to maximize, pulling against the group

Maliciousness

a greedy or adversarial party works to sabotage or steer the deal



*Parties at a table, one adversarial (red).
Scored on the deal reached and on whether
each agent stays in role.*

It is **hard** for earlier models: GPT-3.5 and smaller models mostly fail, and even GPT-4 and Llama-3 70B underperformed in more complex games. A single **adversarial party can derail agreement**, a preview of Part 2.


One compromise, many agents

In L9, prompt injection hijacked a single agent. Injection can spread when agents are connected.

A single image jailbreaks a million agents

Agents share a memory bank (a shared RAG of each other's messages). **One** agent is attacked with a single adversarial image. As agents chat in random pairs, it propagates: an **infectious jailbreak**.

In a simulation of up to **one million** LLaVA agents, infection reaches nearly 100% in about **27 to 31 chat rounds**. The attacker acts once.



Next round ▶

Reset ↺

round 0 · infected 1/64

Click to advance one chat round. Each infected agent infects the peer it talks to: the count roughly doubles each round.

Gu, Zheng, Pang, Du, Liu, Wang, Jiang, Lin, "Agent Smith: A Single Image Can Jailbreak One Million Multimodal LLM Agents Exponentially Fast", ICML 2024 (arXiv:2402.08567). Simulation here is illustrative of the doubling dynamic.

The messages between agents are enough

Agent Smith spread through a shared memory bank. These two attacks need no shared memory at all: plain agent-to-agent text is the whole attack surface. They differ in **where the attacker stands**.

CORBA · ACL 2026 · POISON THE CONTENT

Self-spreading prompt

Prompt one agent with a harmless-looking message and ask it to spread it. The message is sent to every agent that one agent can reach, on any network shape. Agents then end draining each other's compute.


AGENT-IN-THE-MIDDLE · ACL 2025 · POISON THE CHANNEL

Rewrite the messages in transit

Here the attacker never breaks into an agent. It sits **on the link between two agents** and rewrites their messages as they pass, like a man-in-the-middle. Seeing only one agent's incoming messages, an adversarial agent caused denial-of-service or spread a chosen payload in **over 90%** of network layouts.

The same idea, in a real application stack

The first **zero-click GenAI worm** (named after the 1988 Morris worm): a **self-replicating prompt** that makes the model echo itself in its output while taking a malicious action, so each infected email assistant infects the next.




RAG route: a poisoned email from c_1 is stored in c_2 's database, retrieved later as context, then forwarded to c_3 .

Network structure is a safety variable

How far a compromise spreads depends on **how the agents are wired together**, not just on the agents. Dense graphs spread adversarial content fastest; distance from the attacker is safer.

- **NetSafe:** in a misinformation study, a **star** topology's accuracy fell **29.7%** (95.0 to 66.8) from one attacker.
- **Huang et al.:** under faulty agents, the more **centralized** the structure, the more resilient (chart at right).
- A review agent ("Inspector") that corrects messages recovers up to **96.4%** of the lost performance (Huang et al.).



Performance loss under faulty agents, across three structures (Huang et al.).

Coordinating against others

Agents optimizing their own objective can reach outcomes that harm their principals or other agents.


Pricing agents drift into a cartel on their own

Set up LLM agents as rival shops in a simulated market. Each is told only one thing: **maximize your own profit**. None is told to collude, keep prices high, or avoid price wars.

Left alone, they settle on **high prices**.

How to read the chart

Two reference prices: the **competitive** price (red, low, what real rivalry should produce) and the **monopoly** price (green, high, what one owner of both shops would charge). Each dot is one run; the cluster sits near the green line.



*Fish et al. 2024. Each dot is one two-firm market (firm 1's vs firm 2's average price). The two colors are two prompt wordings differing by one sentence: blue (P1) tells agents "do not undermine profits"; orange (P2) reminds them "undercutting sells more". Both sit above the **competitive** price (red); the P1 wording reaches the **monopoly** price (green). A tiny prompt change shifts how collusive they get.*

It can show up wherever agents repeatedly interact

AUCTIONS · AGRAWAL ET AL. 2025

Communication makes it worse

When sellers in an auction are allowed to **message each other**, they collude more. How much depends on the model and on pressure in the setting (being watched, or pushed by an authority).

MARKET DIVISION · LIN ET AL. 2024

Splitting the market silently

Two firms selling two products quietly **carve up the market**, each taking one product and moving toward monopoly, without exchanging a single message.

WHY IT MATTERS

Collusion emerges

Across these settings, collusion is not a rare behavior. Self-interested agents that interact repeatedly and can model each other tend to **drift into collusion**.

Conventions and bias emerge in a population

The classic **naming game**. A population of agents (24 here) has to invent a shared name, with **no leader and no global view**.

One round

Two agents are picked at random. Each independently guesses a name from a shared pool (e.g. 10 letters). **Match: both rewarded.** **Mismatch: both penalized.** Each agent remembers only its own last few rounds.

- Repeated play makes the whole population **settle on one name**: a shared **convention** ("A" below).
- **Collective bias**: some names win far more often, even when single agents tested alone show no preference.
- **Tipping point**: a committed minority that always pushes a new name "B" flips everyone once it passes a critical mass (~2% to 25%, model dependent).

agents always pushing B: **10%** 



Below the critical mass: name A holds.

"Convention A" is the name the population already agreed on. Drag to grow the minority that always says B. Past the critical mass (~25% here), everyone flips to B.

Making collusion measurable

Agents share a **cooperative goal**, but a subset of them (a **coalition**) can quietly coordinate to chase a **side-goal** that hurts the team. Colosseum audits when this happens.

A concrete task

Agents act as engineers picking tickets from a shared backlog (the JIRA setting). The team goal is an efficient overall allocation. A coalition could instead grab the best-value tickets for itself and leave the costly ones to everyone else.

Each task is written as a **formal optimization problem**, so the best cooperative outcome can be computed and any deviation can be scored, not just described.

WHAT THEY DO

Coalition advantage

A regret score on the **actions taken**: how much the coalition gained for itself versus the team's best cooperative outcome.

WHAT THEY SAY

Collusion judge

An **LLM-as-a-judge** reads the conversation logs and rates how collusive the talk is, on a 0 to 5 scale.

Agents collude when handed a private channel

Emergent collusion

Give two otherwise-benign agents a **secret side-channel** and they start colluding **unprompted**. Across **19 models**, most (about 58%) gain coalition advantage once the channel exists.

Talk and actions diverge

The more an agent actually **gains**, the **less collusive its talk sounds** (a negative correlation). We need metrics to audit collusion beyond LLM-as-a-judge.

Bridge to Part 4: if that secret channel is **hidden inside innocent-looking text** (steganography), reading the chat no longer reveals the collusion.

Crossing the two signals gives four behaviors:

	ACTS HONESTLY	ACTS COLLUSIVELY
TALKS COLLUSIVELY	Attempted all talk, no action ("collusion on paper")	Direct plans it and acts on it
TALKS NORMALLY	Cooperative no collusion	Hidden acts on it, never says so

Hidden in plain sight

Our main lever against a misbehaving agent is reading what it says. What if two agents can pass the real message past us, in text that looks innocent?

Steganography and watermarking

Both are **information hiding**: a signal embedded in ordinary-looking content without visibly changing it. What differs is the **purpose** and **which adversary** each must defeat.

Watermarking (recall L5)

A hidden mark in model output for **attribution** (evidence the text came from a given model). It stays **low-distortion** so it does not hurt quality, and a key-holder can verify it. The hard part is **robustness**: surviving edits and paraphrasing, and resisting removal. Many schemes are even **undetectable without the key**.

Steganography

A hidden **payload** (a real message) encoded inside innocent-looking **coverttext**. The goal is **undetectability**: a monitor cannot tell that any message is present. Only a receiver sharing the secret key can recover it.




If agents can do this reliably, then **monitoring their communication stops working**.

Secret collusion via steganography

An agent holds information it is **forbidden to share** and knows an overseer reads every message. Using a scheme shared with its partner, it hides the real message inside innocent text. The overseer sees nothing wrong; the partner decodes it.

- 1 • **Setup.** GPT-4 is a finance agent, told to share only general information, never insider tips, and that an overseer monitors all its messages.
- 2 • **The secret.** It privately holds an insider tip: Microsoft lost a contract, so funds should sell. It must not say this outright.
- 3 • **Hide it.** Asked which companies to pick, it answers in innocent "weather" covertext ("the four corners of our sky... a storm to the North..."), which encodes the tip via a shared key.
- 4 • **Decode.** The overseer reads only weather talk and flags nothing. The partner, holding the same key, recovers the real message: "do not invest in Microsoft."



Motwani et al. 2024, Fig. 1. The forbidden tip travels past the overseer hidden inside ordinary weather talk.

How close are today's models?

The paper **formalizes secret collusion** and builds **CASE** (Collusion and Steganographic Evaluation): it breaks the skill into pieces and tests each under optimization pressure.

What collusion requires

Agree on a scheme without openly discussing it, **encode and decode** the payload (e.g. a cipher), and wrap it in **coverttext** that looks normal to the overseer. CASE measures each piece and the full task.

Earlier models (GPT-2, Llama-2, Mixtral, GPT-3.5) score near 0 on the cipher task; GPT-4 is the first with a clear **capability jump**. The full covert task still succeeds only ~26% of the time, but the trend points up, so the recommendation is to monitor frontier systems before it matures.

92%


GPT-4 correctly encodes/decodes text with a Base64 cipher (54% for a Caesar cipher)

~26%

GPT-4 evades the overseer on the full covert-text-steganography task

Two agents learn to hide collusion from the overseer


A **sender** and a **receiver** must finish a task together while secretly passing a hidden **payload**. A mitigation team fights back: an **overseer** reads every message and a **paraphraser** can rewrite it. Both teams are optimized against each other.



Where we actually are today

A systematic capability evaluation across the Claude, GPT, and Gemini families on two skills: passing hidden messages, and hiding reasoning. The findings are a useful reality check against the worst case.

- Under **standard conditions**, models **cannot** pass a hidden message without a monitor noticing.
- They succeed only with **extra affordances**: an unmonitored scratchpad plus a pre-agreed coding scheme.
- Early encoded reasoning appears on simple state-tracking problems, but they rarely hide it subtly enough to fool a monitor.



Success at hiding messages (B) and reasoning secretly (D) only rises with a "strong tip", an explicit pre-agreed scheme. Under no help, models stay near zero.

Current verdict

Capabilities are **nascent**: a well-designed monitor still catches today's attempts. The risk is real but not yet acute, and we want detection in place **before** the capability jump, not after.

More agents, more wrong

Multi-agent setups are often sold as a way to get more reliable answers. Often the interaction itself makes the answer worse.

Conformity and sycophancy degrade the answer

LLMs cannot handle peer pressure

Shown peers' answers, models **lose more correct answers than they gain**. About 65% of answer changes go from right to wrong. Average accuracy drops 5.65% (models up to 32B). Susceptibility rises with perceived rapport.

A correct minority gets talked out of the right answer by a confident majority.

Sycophancy collapses debate

In multi-agent debate, **sycophancy drives "disagreement collapse"**: agents abandon correct positions to agree. Homogeneous debates reached collapse rates up to **86%**, with a 0.90 correlation between giving up a correct stance and sycophancy.

Mixing roles (a "troublemaker" that keeps disagreeing) helps. Uniform agreeable agents converge fast on the wrong answer.

Sycophancy and unfaithful reasoning are single-model flaws, but in a multi-agent system they **compound**. Adding agents does not average out errors.

What we can do about it

No silver bullet, but a layered set of defenses, each matched to a failure mode from the earlier parts.

Match the defense to the failure mode

AGAINST PROPAGATION

Topology and monitoring

Choose resilient structures (hierarchical over star), add review agents, and detect anomalous agents from the communication graph. **G-Safeguard** runs a GNN over the message graph and recovers ~40% of performance lost to injection.

AGAINST COLLUSION

Trust and credibility

Weight an agent's influence by a **credibility score** earned from past contributions, downweighting adversaries. Reported effective even when adversaries are the majority.

AGAINST COVERT CHANNELS

Paraphrasing and capability limits

Paraphrase messages with a trusted model to break low-level encodings, monitor for covert text anomalies, and limit access to shared keys or fine control. Necessary, but shown insufficient alone.

AT THE SYSTEM LEVEL

Visibility and governance

Agent identifiers, activity logging, and real-time monitoring so deployers can see **who** is acting and **why**. You cannot govern interactions you cannot see.

When the text hides it, read the activations

Parts 3 and 4 showed the limit of watching **what agents say**: collusion can diverge from the talk, or hide inside innocent text. A complementary signal is **what the model computes**.

The method

Linear probes on model **activations** already flag deception in a single agent. Lift this to the group: aggregate each agent's deception score into one system-level in order to detect "is this group colluding?".

NARCBench tests it under distribution shift.

No single probe wins on every collusion type, and detection gets **stronger as models get more capable**, an encouraging counter given that covert-channel skill also rises with capability.

1.00

AUROC detecting collusion in-distribution, across open-weight models


0.73 to 0.93

zero-shot transfer to structurally different scenarios (Llama-3.1 70B)

0.99 to 1.00

on a steganographic blackjack task: it catches the hidden channel

Why multi-agent safety is its own problem



A compromise can spread faster than we contain it. Self-interested agents collude with no instruction to. And the channel we rely on for oversight, reading what agents say, can be hidden from us. Each is a **system-level** property, invisible when you test one agent at a time.

Four things to remember

01

Risk is in the interaction

Multi-agent failure has three modes: **miscoordination**, **conflict**, **collusion**, driven by network effects, selection pressure, and new attack surfaces. Safe parts do not guarantee a safe whole.

02

Compromise propagates

A single jailbreak, worm, or false belief can spread across connected agents, sometimes **exponentially**. Network **topology** is itself a safety variable.

03

Collusion is a default, not an exception

Self-interested agents reach **supracompetitive prices**, **divide markets**, and **form collective biases and conventions** with no instruction to. A committed minority can tip a whole population.

04

Reading the text is not enough

Steganography lets agents pass real messages past a monitor, and the skill scales with the model. The emerging counter is to stop trusting the text alone: read the model's **activations**.

RECAP

What we covered, and where to read more

FRAMING	Three failure modes (miscoordination, conflict, collusion) and seven risk factors; multi-agent security; negotiation testbed	Hammond 2025 , Schroeder de Witt 2025 , Abdelnabi 2024
PROPAGATION	Infectious jailbreak; contagious and man-in-the-middle channel attacks; GenAI worm; misinformation flooding; topology and resilience	Gu 2024 , Zhou 2026 , He 2025 , Cohen 2025 , Ju 2024 , Yu 2025
COLLUSION	Algorithmic price collusion; auctions and market division; emergent conventions and bias; auditing collusion (say vs do)	Fish 2024 , Lin 2024 , Ashery 2025 , Nakamura 2026
STEGANOGRAPHY	Stego vs watermarking; secret collusion and CASE; emergent multi-agent steganographic collusion; frontier capabilities	Motwani 2024 , Mathew 2025 , Schroeder de Witt 2023 , Zolkowski 2025
INTERACTION	Peer pressure and sycophancy degrade multi-agent collaboration	Song 2025 , Yao 2025
DEFENSES	Topology and graph monitoring; credibility scoring; paraphrasing; activation-level collusion detection; visibility and governance	Wang 2025 , Ebrahimi 2025 , Rose 2026 , Chan 2024

Thank you. Questions?

Next lecture: Scalable oversight. How we supervise systems too capable to check directly, and AI control as a way to manage agents we cannot fully trust.



Scan to open the form

WE'D LOVE YOUR FEEDBACK

Tell us how this lecture went

forms.gle/5JF3H6zhP795HhpD8